

Thai Poetry Rhyme Verification using a Hybrid Rule-Based Approach for Education and GenAI

1st Warit Yuvaniyama
School of Information, Computer,
and Communication Technology
Sirindhorn International Institute
of Technology, Thammasat University
Pathum Thani 12120, Thailand
<https://orcid.org/0009-0007-1630-4286>

2nd Athit Phinyachok
School of Information, Computer,
and Communication Technology
Sirindhorn International Institute
of Technology, Thammasat University
Pathum Thani 12120, Thailand
<https://orcid.org/0009-0000-3783-8104>

3rd Gunn Sanguankittipun
School of Information, Computer,
and Communication Technology
Sirindhorn International Institute
of Technology, Thammasat University
Pathum Thani 12120, Thailand
<https://orcid.org/0009-0005-8812-4992>

4th Leo Colombo
School of Information, Computer,
and Communication Technology
Sirindhorn International Institute
of Technology, Thammasat University
Pathum Thani 12120, Thailand
<https://orcid.org/0009-0004-9133-5398>

5th Prachya Boonkwan*
School of Information, Computer,
and Communication Technology
Sirindhorn International Institute
of Technology, Thammasat University
Pathum Thani 12120, Thailand
<https://orcid.org/0000-0003-1405-1071>*

Abstract—In Thai Klon-Paed (กลอนแปด) poetry (the eight-syllable verse form) each four-line stanza must satisfy an interlocking rhyming pattern where each rhyming syllables must share both a vowel sound (sara สระ) and a final-consonant class (mattara มาตราตัวสะกด). Because Thai is written without word boundaries and its orthography alters form or conceals implicit vowels, silenced letters, and consonant clusters. Automatic verification of rhymes is non-trivial, and neither classical rule-based checkers nor grapheme-to-phoneme (G2P) romanization checkers deal with the resulting edge cases reliably. We present a hybrid rule-based verification approach. First, we improve the KhaveeVerifier algorithm—vowel analysis, final-consonant-class analysis, the rhyme test, karun silencing, and true-final detection—and ship it in PyThaiNLP 5.3.5 (Model B). Second, we introduce Klonpad (Model D), which augments these rules with a gold-standard G2P override dictionary of 4,292 entries and a fallback segmenter, recovering phonetic blind spots of the base oracle. We evaluate five checker configurations on a gold corpus of 36,475 classical stanzas (145,709 rhyme checks) and on a targeted augmentation of 11,623 instances that includes 10,000 adversarial negatives and 423 oracle-blind rhymes. The best configuration reaches 88.5% gold recall and 86.9% F1 on the augmentation while maintaining 100% precision on curated negatives and recovering 95.7% of the oracle-blind rhymes. The verifier serves both as an explainable, zero-hallucination educational tool and as a deterministic reward model for reinforcement-learning-based poetry generation.

Index Terms—Thai NLP, Thai poetry, Klon-Paed, rhyme verification, rule-based systems, grapheme-to-phoneme, reinforcement learning, educational technology

I. INTRODUCTION

Thai orthography does not map straightforwardly onto pronunciation, and this gap becomes especially apparent in the classical register used by Thai poetry. The script is written without explicit word boundaries; vowels are frequently im-

PLICIT; consonants or entire clusters can be silenced by the thanthakhat (ทัณฑฆาต) (karun) diacritic; and irregular words whose pronunciations do not follow the standard phonetic rules are common. Complex phonological structures, such as consonant blending (*kham khuap klam* คำควบกล้ำ) and leading consonants (*aksorn nam* อักษรนำ), further obscure syllable boundaries, vowel quality, and tone in an already unsegmented writing system. As a result, syllabification, pronunciation, and tone assignment are genuinely difficult tasks for both human learners and computational systems.

Klon-Paed, the eight-syllable Thai verse form, depends on exactly these phonetic properties. Each stanza (*bot* บท) consists of four sets of seven-to-nine-syllable lines (*wak* วรรค). A valid rhyme between two syllables requires that they share both the vowel sound (sara สระ) and the final-consonant class (mattara มาตราตัวสะกด). Four canonical rhyme rules must hold within (r1, r2, and r3) and across stanzas (rX) (Section II), and the final syllable of every wak participates in the scheme. A single silent letter—for example the silent *r* of *phhet* (เพชร) (‘diamond’)—or a misread vowel length is enough to break a rhyme. These constraints are strict and deterministic, which is precisely what makes them hard for statistical systems.

Two communities need a reliable automatic verifier for Klon-Paed. In *education*, teachers and students need feedback that is deterministic, explainable, and free of hallucination: an LLM that confidently invents a rhyme that does not exist is worse than useless in a classroom. In *NLP research*, training a language model to compose Klon-Paed, for instance through reinforcement learning (RL), requires a programmatic reward signal that scores rhyme correctness exactly. Because large language models (LLMs) are probabilistic engines, they are

intrinsically unreliable at strict constraints such as syllable counts and rhymes; a rule-based verifier supplies the ground truth that such generative pipelines require.


Existing verification systems fall short of this standard. The historical KhaveeVerifier in PyThaiNLP 5.0.1 (Model A) relies on dictionary subword tokenisation, omits one of the four canonical rhyme rules, and reaches only 73.4% recall on our gold corpus. Grapheme-to-phoneme (G2P) romanisation-based checkers, such as the `word_check` grader of the KlonSuphap-LM generator [11] (Model C), collapse long/short vowel distinctions and produce thousands of false positives on adversarial negatives. Neither family handles the phonetic blind spots of Thai orthography in a principled way. In this paper we address these limitations with a hybrid rule-based approach and a rigorous, adversarial evaluation. Our contributions are:

- **An improved rule-based verifier (Model B).** We rework the vowel analysis (`check_sara`), the final-consonant-class analysis (`check_marttra`), the rhyme test (`is_sumpus`), karun silencing, and true-final detection, and implement all four canonical rhyme rules. The result is merged into PyThaiNLP 5.3.5 as the default KhaveeVerifier.
- **A dictionary-enhanced variant (Klonpad, Model D).** We augment the rules with a gold-standard 4,292-entry G2P override dictionary and a fallback segmenter, recovering the oracle’s blind spots.
- **A large evaluation corpus.** 36,475 gold stanzas (145,709 rhyme checks) from five classical Thai works, plus 11,623 targeted augmentation instances—10,000 oracle-verified negatives, 1,200 tricky positives, and 423 oracle-blind positives—each audited in independent review passes.
- **A systematic comparison.** Five configurations (A, B, C, D_w2p, D_ssg) are compared on precision, recall, F1, coverage, and confidence intervals, with per-rule and per-operator error analyses.
- **Two deployments.** An explainable educational web tool and a deterministic reward environment for RL-based Klon-Paed generation.

II. BACKGROUND AND RELATED WORK

A. The Linguistic Rules of Klon-Paed

A Klon-Paed stanza has four waks of eight syllables each. Thai syllables are written without delimiters and a word may consist of one or more syllables, so the metre must be recovered by syllabification rather than by counting words. Within this skeleton the verse obeys a fixed rhyme scheme. Two syllables rhyme (*sumpus* สัมผัส) when they share the same vowel sound (*sara* สระ) and the same final-consonant class (*mattrā* มาตราตัวสะกด), as defined by the Royal Institute of Thailand; tone is not part of the rhyme. The *mattrā* classes (*mae* แม่) —ka, kong, kok, kot, kon, kop, kom, koey, and koewa— group the final consonants of Thai syllables by their phonetic outcome (the velar, dental, labial, nasal, semivowel, and w finals). The four waks are the (*sadap* สดับ), (*rap* ราบ), (*rong* ร้อง),



ถึงบางพุดพุดดีเป็นศรีศักดิ์ มีคนรักสรรถ้อยร้อยจิต แม้นพุดชั่วตัวตายทำลายมิตร จะชอบผิดในมนุษย์เพราะพุดจา

Fig. 1. The example stanza (Section 2.1) with its rhyme positions shaded and linked: r1 (*sadap*→*rap*), r2 (*rap*→*rong*), r3 (*rong*→*song*); rX (not shown) links wak 4 of one stanza to wak 2 of the next. Syllables are romanised; the original Thai verse is shown below in the XeLaTeX build.

and (*song* ส่ง). Four canonical rhyme rules are checked in this work, following the *klon* canon and the implementation of the PyThaiNLP KhaveeVerifier:

- r1 (*sadap* → *rap*): the final syllable of wak 1 must rhyme with one of the first five syllables of wak 2 (syllables 3 and 5 are obligatory, 1, 2, and 4 are permissible);
- r2 (*rap* → *rong*): the final syllable of wak 2 must rhyme with the final syllable of wak 3;
- r3 (*rong* → *song*): the final syllable of wak 3 must rhyme with one of the first five syllables of wak 4;
- rX (inter-stanza): the final syllable of wak 4 of the previous stanza must rhyme with the final syllable of wak 2 of the current stanza.

As a concrete example, the well-known stanza of the *Phukao Thong* (นิราศภูเขาทอง) travel poem by Sunthorn Phu satisfies the intra-stanza rules r1–r3:

thueng bang phut phut dii pen si sak
mii khon rak rot thoi a-roi cit
maen phut chua tua taai thamlai mit
ca chop phit nai manut phro phut caa
ถึงบางพุดพุดดีเป็นศรีศักดิ์ มีคนรักสรรถ้อยร้อยจิต
แม้นพุดชั่วตัวตายทำลายมิตร
จะชอบผิดในมนุษย์เพราะพุดจา

Figure 1 shows the intra-stanza structure.

Three orthographic phenomena make these rules hard to apply automatically. First, compound, transformed, and reduced vowels (*sara prasom* สระประสม), (*sara plianrup* สระเปลี่ยนรูป), (*sara lotrup* สระลดรูป) let several spellings map to one vowel sound. Second, the *thanthakhat* (karun) diacritic silences a letter or a whole cluster, so the written form contains letters that are not pronounced. Third, a trailing *y*, *l*, or *w* may be a genuine final consonant or merely part of a vowel digraph or an initial cluster. Any verifier must resolve all three before it can compare vowel and final class.

B. Thai Text-Processing Foundations

Two building blocks underlie all Thai NLP systems: segmentation and grapheme-to-phoneme conversion. Because Thai has no word delimiters, segmentation has received sustained attention: dictionary-based maximum matching (the `newmm` engine of PyThaiNLP [1]), neural segmenters such

as DeepCut [5] and AttaCut [3], syllable-based neural segmentation [4], and the SSG syllable segmenter shipped in PyThaiNLP 5.x, which the present work uses throughout. G2P conversion maps surface orthography to pronunciation; Thai G2P has been addressed with rule-based systems such as TLTK [6], example-based methods [7], CRF-based joint-sequence models [8], and two-stage sequence models [9]. G2P is directly relevant to rhyme verification because rhyme is a property of *sound*, not of spelling; Model C in this study performs its comparison in the romanised space produced by the TLTK G2P model.

C. Computational Poetry Analysis and Generation

Prior work on Thai poetry in NLP has concentrated on analysis and generation. On the analysis side, machine-learning studies of Sunthorn Phu’s verse have extracted sound patterns—rhyme, assonance, and alliteration—using internal-rhyme rules [10]. On the generation side, the dominant line of work fine-tunes language models on corpora of classical verse: the Phra Aphai Mani LM [12] and its successor KlonSuphap-LM [11] fine-tune a GPT-2 Thai base model, then use *rhyme tagging* (annotating the tokens that must rhyme), attention masking, and finally reinforcement learning in which a rule-based *klon grader* supplies the reward. That grader is the `word_check` module evaluated here as Model C, making our evaluation directly relevant to the RL step of that pipeline. More broadly, RLHF has established reward models as essential for aligning LLMs to follow strict requirements [2], and grammar-constrained decoding has been proposed to enforce structural constraints at inference time [14]. For classical Thai specifically, recent benchmarks report that even strong LLMs struggle with archaic vocabulary and verse structure [13], which motivates dedicated deterministic tools.

D. Positioning of This Work

Compared with this landscape, our contribution is the first systematic, adversarially augmented evaluation of Thai rhyme verification as opposed to generation. The historical rule-based checker (Model A) has never been stress-tested against curated edge cases; the G2P-based grader of KlonSuphap-LM (Model C) has not been validated for precision; and no open verifier covers all four rhyme rules while providing dictionary support for irregular pronunciations. We fill this gap with improved rule-based verification (Model B), a dictionary-enhanced variant (Model D), and a large gold-plus-augmentation evaluation that quantifies precision and recall, including a novel *oracle-blind* probe that holds the verifier itself accountable for rhymes its own oracle cannot see.

III. METHODOLOGY: HYBRID VERIFICATION ARCHITECTURE

A. Core Rule-Based Verification (Model B)

Model B is the `KhaveeVerifier` of PyThaiNLP 5.3.5. It verifies a poem in three stages. (1) Each wak is tokenised into syllables with the SSG segmenter. (2) For each canonical rule the relevant syllables are extracted: the final syllable of every wak, and the first five syllables of waks 2 and 4. (3) The rhyme

Algorithm 1 Rhyme verification of a Klon-Paed poem (`check_klon`).

Require: poem T with waks separated by whitespace; type $k \in \{4, 8\}$

Ensure: list of rhyme errors (empty $\Leftrightarrow$ poem is valid)

```

1:  $W \leftarrow$  split  $T$  into waks, grouped into stanzas of four
2: for each stanza  $(w_1, w_2, w_3, w_4)$  do
3:   for  $i \leftarrow 1$  to 4 do
4:      $S_i \leftarrow$  ssg( $w_i$ ) {syllable tokens of wak  $i$ }
5:   end for
6:   for each rule  $\rho \in \{r1, r2, r3, rX\}$  do
7:      $(a, B_\rho) \leftarrow$  source syllable  $a$  and target set  $B_\rho$  of rule  $\rho$ 
8:     if  $\nexists b \in B_\rho$  with is_sumpus( $a, b$ ) then
9:       record an error for rule  $\rho$ 
10:    end if
11:  end for
12: end for
13: return the list of errors

```

test `is_sumpus` decides whether two syllables rhyme; a rule holds if the test succeeds for at least one permitted target. The eight-syllable metre is additionally reported. Algorithm 1 summarises this flow.

The accuracy of the verifier rests on three functions. `check_sara` returns the canonical vowel of a syllable. It handles simple vowels, compound vowels (*sara prasom*), transformed and reduced vowels (*sara plianrup/lotrup*), the special vowels (*sara koen* สระเกิน): *am*, *ai*, *ao*, Pali/Sanskrit words with silent terminal vowels, and words whose karun-stripped form changes the vowel. `check_marttra` returns the Royal-Institute final-consonant class of a syllable, strictly by orthography. It resolves karun first, strips a silent *r* in the *tr/chr/pr* clusters of loanwords (e.g., *but*, *net*, *phhet* (เพชร)), handles single-character words, and classifies the *sara koen* orthographically as *mae ka*; phonetic rhyming for these is restored later in `is_sumpus`. Finally, `is_sumpus` applies phonetic normalisation for the *sara koen*: *a+koey* $\rightarrow$ *ai+ka* (so *chai* rhymes with *sai*) and *a/am+kom* $\rightarrow$ *am+ka* (so *cham* rhymes with *kam*). It returns true iff the normalised sara and mattra of the two syllables are equal.

Two further mechanisms account for most of the improvement over the 5.0.1 baseline. `handle_karun_sound_silence` strips characters silenced by the thanthakhat (e.g., *ohm* (โหม) $\rightarrow$ *o*), so that *ohm* and *o* are judged to rhyme. `_is_true_final` decides whether a trailing *y*, *l*, or *w* is a genuine final consonant rather than part of a vowel digraph or an initial cluster: the *y* of *thai* is silent, the *l* of *plae* belongs to the initial cluster, but the *l* of *chon* is a true final.

Table I summarises the principal changes of the 5.3.5 verifier relative to the 5.0.1 baseline, each with a concrete example; they were developed and merged in the PyThaiNLP project [1]. The fixes target precisely the failure families probed by the augmentation (Section V): karun and *ru* traps, short/long

TABLE I
PRINCIPAL CHANGES OF THE 5.3.5 KHAVEEVERIFIER OVER THE 5.0.1
BASELINE.

Aspect	5.0.1 → 5.3.5 (example)
Segmentation	dict subword → SSG syllable segmentation
Karun (thanthakhat)	ignored → silenced (ohm→o)
True finals y/l/w	naive → cluster-aware (thai, klai→mae ka)
Special-vowel norm.	one-sided → both-sided (cham~kam)
Compound/ reduced vowels	misread → compound-vowel classifier, short-vowel mark, and ua diphthong (chet→short e; tua~chua)
ru	one sound → three sounds (roek, rit, rue-du)
r3 rule	absent → implemented (all four rules)
Silent r	misread → stripped in check_matttra (but, phet)
Single-char words	crash → guarded (na, tha→a)

and compound-vowel swaps, and the one-sided special-vowel normalisation that our C9 operator exploits. Together they raise per-rule gold recall from 86–89% (5.0.1) to 95–97% (5.3.5) and eliminate false positives on the curated negatives.

B. Dictionary-Enhanced Verification (Klonpad, Model D)

Model B is linguistically principled, but it inherits a blind spot from its own design: `check_sara` returns only the first vowel of a token, so a multi-syllable word that SSG keeps as one token loses the vowel of its final (rhyming) syllable, and the silent-*r* loanwords are stripped only in a whitelisted subset of clusters. Klonpad (Model D) addresses this by layering a gold-standard pronunciation dictionary on top of Model B’s rules.

The dictionary, `POETRY_OVERRIDES`, contains 4,292 hand-verified entries mapping a surface form to its syllable sequence (e.g., *phhet* (พหฺ) → *phet*; *sattruu* (ศัฏฐ) → *sat-truu*; *kasatrii* (กษัตริ) → *ka-sat-trii*). It was built from the frequent vocabulary of the Phra Aphai Mani corpus with manual supervision, and every entry is checked by a four-guard pipeline that rejects pronunciations hallucinated by the Word2Phrase (w2p) neural G2P model; w2p-derived candidates are kept only as labelled hints and are never trusted on their own. At inference time each syllable is first looked up in the override dictionary; if absent, a fallback segmenter supplies the pronunciation. We report two configurations: `D_w2p`, which falls back to the w2p neural G2P model, and `D_ssg`, which falls back to plain SSG segmentation without w2p. As Section V shows, the w2p fallback introduces pronunciation hallucinations on short or compound words, and `D_ssg` is the more robust configuration.

C. Baseline Systems

Model A is the original `KhateeVerifier` of `PyThaiNLP` 5.0.1, instrumented verbatim from the vendored source: dictionary-based `subword_tokenize`, no `r3` rule, no `karun` handling, and no true-final logic. Model C is

TABLE II
GOLD CORPUS STATISTICS (ALL-POSITIVE).

Work	Files	Stanzas	Rhyme checks
Suphasaet Son Ying	1	200	799
Khobut	14	1,303	5,198
Khun Chang Khun Phaen	43	10,543	42,129
Phra Aphai Mani	132	24,342	97,236
Phukao Thong	1	87	347
Total	191	36,475	145,709

the `word_check` rhyme grader of the `KlonSuphap-LM` project [11]: each wak is romanised with the `TLTK G2P` model and the romanised vowels and finals are compared. Its two documented weaknesses are (i) a `replace_long_short` step that collapses long and short vowels, and (ii) a `matttra` comparison that ignores the final consonant; both inflate recall at the cost of precision, and `tltk` failures reduce coverage.

IV. DATASET CONSTRUCTION AND AUGMENTATION

A. Gold-Standard Corpus Compilation

We collected five classical Thai narrative poems from the Vajirayana digital library [15]: *Khobut* (โคบุตร), *Khun Chang Khun Phaen* (ขุนช้างขุนแผน), *Phra Aphai Mani* (พระอภัยมณี), *Phukao Thong* (ภูเขาทอง), and *Suphasaet Son Ying* (สุภาษิต สอนหญิง). A cleaning pipeline normalised the text and split it into stanzas of four waks, discarding only chapter-opening fragments and scribal *o* marks. Crucially, waks with ten or more syllables were retained rather than dropped: removing them silently severs the inter-stanza rhyme chains and made an early version of the corpus appear to violate `rX` at a 72% rate. On the complete corpus `rX` holds at 96–97% for the strongest checkers, confirming that the fix was correct. The final corpus comprises 191 files, 36,475 stanzas, 145,900 waks, and 145,709 gold rhyme checks (Table II); the gold is all-positive, because the classical poems are presumed to rhyme. An alternative public corpus covering similar material is the Thai Literature Corpora [16].

B. Targeted Data Augmentation

Recall over an all-positive corpus cannot measure precision, and easy positives cannot expose blind spots. We therefore augmented the gold with adversarial instances generated by an automated pipeline. Every instance is *standalone*: it is derived from a real stanza but only the target rule is changed, so exactly one rule is exercised while the others remain intact. Every negative is verified by the 5.3.5 oracle (`is_sumpus`) to be a genuine non-rhyme, and the whole set was audited in four independent passes in which reviewers spelled out the `sara` and `matttra` of every candidate (with vowel length made explicit) and double-checked pronunciations against the IPA. The audit removed 56 words from the candidate pool that the oracle misreads—`karan` clusters it fails to silence, and multi-syllable Pali words SSG keeps as single tokens.

TABLE III
AUGMENTATION COMPOSITION (11,623 INSTANCES).

Operator	Effect tested	Count
C6 random	different sara and matta (baseline)	300
C0 same-matta	same matta, different sara	500
C1 same-vowel	same sara, different matta	500
C3 short-long	short↔long vowel swap	3,100
C4 old-disagree	r/l/w finals, 5.0.1 disagrees	1,200
C9 old-accept	5.0.1 accepts, 5.3.5 rejects	2,800
C5 lead-head	leading h/o words	600
C2 trap	karun / ru / cluster traps	1,000
HP tricky	curated hard positives	1,200
HP oracle-blind	rhymes the oracle cannot see	423

a) *Negatives (precision probes)*.: The 10,000 negatives are drawn from eight corruption operators (Table III): trivial different-vowel-and-matta swaps (C6), same-matta swaps (C0), same-vowel swaps (C1), short/long vowel swaps (C3), liquid-final swaps (C4), a systematic trap in which the 5.0.1 checker reads the candidate as rhyming with the target while 5.3.5 does not (C9), leading-consonant words (C5), and karun/ru (ရ)/cluster traps (C2). C3 and C9 dominate the mix because they are the operators that best discriminate the checkers.

b) *Positives (recall probes)*.: The 1,200 *tricky* positives swap the target syllable with a rhyming candidate that exercises normalisation (*sara plianrup/lotrup* *สระเปลี่ยนรูป/ลดรูป*), *ru* (ရ), *karun*, and true-final cases; all are oracle-confirmed rhymes. A further 423 *oracle-blind* positives are genuine rhymes that the oracle itself cannot verify and therefore labels as non-rhymes: the silent-*r phhet* (ဖုတ်) (true *eh+kot*, which the oracle reads with the long *ae* because the dead final carries no short-vowel mark), and the first-sara words *sattruu* (နတ်), *kasatrii* (နတ်), and *kasatriiya* (နတ်) (SSG keeps each as one token, so *check_sara* hears only the first vowel). For these instances the gold label is the *linguistic* truth, not the oracle verdict—a checker that hears the real rhyme receives credit, while the oracle itself (Model B) is charged a false negative. This is the quantitative statement of “the oracle that is wrong loses points.”

V. EVALUATION AND RESULTS

A. Experimental Setup

We report precision, recall, and F1 with Wilson 95% confidence intervals, coverage (the share of units evaluated), and a conservative drop-as-fail variant for Model C. Because every augmentation instance is standalone and tests exactly one rule, the per-rule prediction is the clean metric; a stanza-level aggregate over the augmentation is only well-defined for Model B (which is the oracle). Model A has no r3 and is reported as N/A there. All checkers were instrumented from verbatim vendored sources (the 5.0.1 and 5.3.5 cores), with Model C run under Python 3.12 through a persistent pool of ten *tltk* workers (the full augmentation pass takes about 40 minutes; A/B/D are parallelised and finish in under a

TABLE IV
GOLD RECALL (STANZA LEVEL AND PER RULE), 36,475 STANZAS. “—” = N/A (NO R3).

System	Stanza	r1	r2	r3	rX
A (5.0.1)	73.4	86.3	88.8	—	88.7
B (5.3.5)	87.5	94.7	96.5	95.5	96.7
D_w2p	87.4	95.4	96.3	94.9	96.5
D_ssg	88.5	95.9	96.7	95.2	97.0
C (Kongfha)	84.8 [†]	93.2	96.5	94.1	96.6

[†] coverage 95.7%; per-rule figures on the evaluated subset.

minute). Every number in this section is recomputed from the committed result files.

B. Gold-Corpus Recall

Table IV reports stanza-level and per-rule recall over the 36,475 gold stanzas. The proposed family dominates the baseline: D_ssg reaches 88.5% stanza recall and B 87.5%, versus 73.4% for Model A. Model C reaches 84.8% but evaluates only 95.7% of the corpus (*tltk* failures concentrate in the archaic Khun Chang Khun Phaen); its figures use the coverage-corrected denominator.

C. Augmentation: Precision, Recall, and F1

Table V summarises the augmentation-only results (pooled per rule) and Table VI breaks them down per rule. Three results stand out.

First, Model B is a perfect oracle on the negatives: it makes zero false positives across all 10,000 curated negatives (100% precision by construction, since the negatives were oracle-verified). Its recall is limited only by the oracle-blind positives it cannot see (73.9%).

Second, Model D recovers most of those blind spots: D_w2p reaches 92.1% recall and D_ssg 87.0%, with precision 83.4% and 86.8% respectively, giving the best F1 on the augmentation (87.5% for D_w2p). The oracle-blind column of Table VII shows the mechanism: D_w2p hears 405 of the 423 oracle-blind rhymes (95.7%), D_ssg 288 (68.1%), whereas B hears none.

Third, the baselines invert the precision/recall trade-off. Model C is recall-strong (87.1% on the augmentation, 78.7% even on oracle-blind rhymes, because its G2P romanisation hears the real sound) but makes 3,107 false positives—2,702 of them on the short/long vowel swap (C3), which its *replace_long_short* collapses—for an augmentation precision of only 30.4%. Model A is weak on both sides (28.4% precision, 62.7% recall), with 1,774 of its 1,950 false positives coming from the C9 operator (words the 5.0.1 *is_sumpus* reads with the same sara/matta as the target while 5.3.5 splits them).

D. Oracle-Blind Probe and Error Analysis

Table VII reports the oracle-blind probe together with the combined false-negative accounting (tricky + oracle-blind positives) and the false-positive breakdown by the two most informative operators. The combined FN-pos column is the

TABLE V

AUGMENTATION-ONLY RESULTS (POOLED PER RULE, 11,623 INSTANCES). FP = FALSE POSITIVES ON 10,000 NEGATIVES; FN-pos = COMBINED FALSE NEGATIVES ON 1,623 POSITIVES.

System	P	R	F1	FP	FN-pos
A (5.0.1)	28.4	62.7	39.1	1,950	848
B (5.3.5)	100.0	73.9	85.0	0	423
D_w2p	83.4	92.1	87.5	298	128
D_ssg	86.8	87.0	86.9	215	211
C (Kongfha)	30.4	87.1	45.1	3,107	266

TABLE VI

AUGMENTATION-ONLY PER-RULE PRECISION/RECALL (%). “—” = N/A (NO R3).

System	r1		r2		r3		rX	
	P	R	P	R	P	R	P	R
A	31.8	52.7	27.8	65.7	—	—	25.7	76.9
B	100	54.5	100	77.5	100	78.9	100	98.0
D_w2p	80.3	91.8	94.8	93.3	71.8	92.4	94.6	90.8
D_ssg	83.1	82.4	98.5	85.5	75.3	90.5	98.3	92.8
C	35.2	83.0	33.7	96.0	24.4	78.2	27.3	93.8

headline number: B carries 423 false negatives (all oracle-blind), A carries 848, and D carries only 128 (D_w2p) or 211 (D_ssg). The per-operator FP columns show that C’s precision collapse is concentrated on the short/long vowel swap (C3: 2,702 of 3,107 FPs) and A’s on the old-accept trap (C9: 1,774 of 1,950).

E. Discussion

The experiment supports four findings. (1) **The proposed rules are precise.** B never false-accepts on oracle-verified negatives, and its precision transfers to the merged gold+augmentation setting, where B, D_w2p, and D_ssg all exceed 99.5% precision while A and C stay at 95% and 92%. (2) **The override dictionary recovers the oracle’s blind spots.** The D_w2p-versus-D_ssg gap on oracle-blind recall (95.7% versus 68.1%) is a clean, mechanistic demonstration of the pronunciation-knowledge contribution: the curated overrides and w2p both read *phhet* as /phet/ and split *sattruu* and *kasatrii* into their true syllables. (3) **C trades precision for recall on exactly the vowel-length distinction** that the rule-based family enforces; its G2P romanisation is useful for hearing real rhymes (78.7% oracle-blind) but its length collapse makes it unusable as a precision grader. (4) **The oracle itself has a blind spot that only the oracle-blind probe exposes.** A purely oracle-supervised evaluation would have reported B as flawless, hiding the deduction that our probe applies.

Overall, D_ssg is the best configuration on the gold corpus (88.5% recall) and D_w2p the best on the augmentation (87.5% F1). On the merged gold+augmentation set, the F1 favours D_ssg (93.4%), followed by D_w2p (92.9%) and B (92.8%), with C (88.0%) and A (82.3%) behind; the merged precision is gold-dominated (all checkers $\geq 92\%$), so the precision that discriminates the systems lives in the augmentation-only tables.

TABLE VII

ORACLE-BLIND PROBE AND ERROR ANALYSIS. FN-pos = COMBINED TRICKY + ORACLE-BLIND FALSE NEGATIVES.

System	recall (%)		FN-pos		FP	
	tricky	oracle-blind	total	=ob	total	C3/C9
A	57.0	21.5	848	332	1,950	162/1,774
B	100	0.0	423	423	0	0/0
D_w2p	90.8	95.7	128	18	298	135/103
D_ssg	93.7	68.1	211	135	215	119/48
C	85.3	78.7	266	90	3,107	2,702/267

VI. APPLICATIONS AND FUTURE WORK

A. Educational Impact

The verifier is deployed as an interactive web tool for teachers and students of Klon 4/8. The interface accepts a poem, segments it into syllables, and displays each wak with its rhyme positions highlighted and the sara and mattrā of every syllable spelled out. Errors are reported deterministically and explained in terms of the four rhyme rules, so a student can see exactly why a line fails (for example, a short/long vowel mismatch or a misread karun). Because the system is rule-based, its feedback is reproducible and free of hallucination—properties that general-purpose LLMs cannot guarantee. Future work includes classroom evaluations of the tool and pronunciation guidance for irregular words.

B. Foundation for Generative AI

The same determinism makes the verifier an ideal reward model for training small, locally run LLMs to compose Klon-Paed. Reinforcement learning requires a dense, consistent reward signal; the verifier provides one that is exact, immediate, and free of the noise of a learned reward model. This mirrors the RL step of the KlonSuphap-LM pipeline [11], but with a grader (our Model B or D) that our evaluation shows to be substantially more precise than the G2P-based grader used there (Model C), particularly on vowel length. Future work includes scaling the override dictionary to more corpora, adding context-sensitive pronunciations for heteronyms, and integrating the verifier into a decoding-time constrained sampler for poetry generation [14].

VII. CONCLUSION

We presented a hybrid rule-based approach to Thai Klon-Paed rhyme verification. An improved rule-based verifier (Model B), now the default *KhaveeVerifier* of *PyThaiNLP* 5.3.5, implements all four canonical rhyme rules with robust handling of karun, true finals, and sara normalisation, and reaches 100% precision on curated negatives. A dictionary-enhanced variant (Klonpad, Model D) adds a gold-standard 4,292-entry G2P override dictionary and recovers 95.7% of the oracle’s own blind spots. Evaluated on 36,475 gold stanzas and 11,623 adversarial instances against two external baselines, the proposed family raises stanza recall from 73.4% (*PyThaiNLP* 5.0.1) to 88.5% (D_ssg) while maintaining near-perfect precision, and exposes—quantitatively—the blind spots of both the

G2P-romanisation checker and the rule-based oracle itself. The verifier is released as an explainable educational tool and as a deterministic reward environment for reinforcement-learning poetry generation, and its integration into PyThaiNLP makes the improvement available to the entire Thai NLP ecosystem.

ACKNOWLEDGMENT

We thank the Vajirayana Digital Library [15] for the literary datasets, the PyThaiNLP maintainers for reviewing and merging the `KhaveeVerifier` improvements, and the developers of TLTK and the KlonSuphap-LM project for the reference implementations used as baselines.

REFERENCES

- [1] PyThaiNLP Contributors, “PyThaiNLP: Thai natural language processing in Python,” 2024. [Online]. Available: <https://github.com/PyThaiNLP/pythainlp>
- [2] L. Ouyang, J. Wu, X. Jiang, et al., “Training language models to follow instructions with human feedback,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [3] P. Chormai, P. Prasertsom, and A. T. Rutherford, “AttaCut: A fast and accurate neural Thai word segmenter,” *arXiv preprint arXiv:1911.07056*, 2019.
- [4] P. Chormai, P. Prasertsom, J. Cheevaprawatdomrong, and A. T. Rutherford, “Syllable-based neural Thai word segmentation,” in *Proc. 28th International Conference on Computational Linguistics (COLING)*, 2020, pp. 4617–4628.
- [5] T. Kittinaradorn, et al., “DeepCut: A Thai word tokenization library using deep neural network,” 2019. [Online]. Available: <https://github.com/rkcosmos/deepcut>
- [6] W. Aroonmanakun, “TLTK: Thai language toolkit,” [Online]. Available: <https://pypi.org/project/tltk/>
- [7] P. Charoenpornasawat and T. Schultz, “Example-based grapheme-to-phoneme conversion for Thai,” in *Proc. Interspeech*, 2006.
- [8] S. Saychum, S. Kongyoung, A. Rugchatjaroen, P. Chootrakool, S. Kasuriya, and C. Wutiwiwatchai, “Efficient Thai grapheme-to-phoneme conversion using CRF-based joint sequence modeling,” in *Proc. Interspeech*, 2016.
- [9] A. Rugchatjaroen, S. Saychum, S. Kongyoung, P. Chootrakool, S. Kasuriya, and C. Wutiwiwatchai, “Efficient two-stage processing for joint sequence model-based Thai grapheme-to-phoneme conversion,” *Speech Communication*, 2019.
- [10] P. Suksanguan, S. Waijanya, and N. Promrit, “The extraction of beautiful sound patterns from Sunthorn Phu’s poem using machine learning technique and internal rhyme rule,” 2021.
- [11] K. Yingyingseree, “KlonSuphap-LM: Thai Klon-Paed generation with GPT-2,” *Hugging Face model and GitHub repository*, 2024. [Online]. Available: <https://huggingface.co/Kongfha/KlonSuphap-LM>
- [12] K. Yingyingseree, “PhraAphaiManee-LM,” *Hugging Face model*, 2023. [Online]. Available: <https://huggingface.co/Kongfha/PhraAphaiManee-LM>
- [13] T. Triyason, “Lilit to Ramakien: A five-work benchmark for LLM understanding of classical Thai literature,” in *Proc. International Conference on Advances in Information Technology, ACM*, 2026.
- [14] S. Geng, M. Josifoski, M. Peyrard, and R. West, “Grammar-constrained decoding for structured NLP tasks without fine-tuning,” in *Proc. Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [15] Vajirayana Digital Library. [Online]. Available: <https://vajirayana.org>
- [16] A. Rutherford, “Thai literature corpora (TLC),” [Online]. Available: <https://attapol.github.io/tlc.html>